

Data Analysis and Text Mining

Project

The goal of this project is to process a text dataset using data mining and machine learning methods and to extract knowledge from it. Prepare a report for this Project.

In this Project you will use *20 Newsgroup dataset*, but you are free to use any text data (UCI datasets repository, kaggle, data.gouv.fr, ...)

The project should contain at least the following 4 parts:

1. Text processing and Transformation
2. Analysis of the dataset
3. Machine learning and prediction
4. Visualization of the prediction results
5. Theoretical formalism

1. Text Processing and Transformation

For this part, you should use scikit-learn and you can follow the tutorial:

https://scikit-learn.org/stable/tutorial/text_analytics/working_with_text_data.html#tutorial-setup

A. Extracting features from text files

In order to perform machine learning on text documents, we first need to turn the text content into numerical feature vectors using:

- Bags of words:

The most intuitive way to transform the text into a vector is to use a bags of words representation:

1. Assign a fixed integer id to each word occurring in any document of the training set (for instance by building a dictionary from words to integer indices).
2. For each document i , count the number of occurrences of each word w and store it in $X[i, j]$ as the value of feature j where j is the index of word w in the dictionary.

- Tokenizing text with scikit-learn

Text preprocessing, tokenizing and filtering of stopwords are all included in CountVectorizer, which builds a dictionary of features and transforms documents to feature vectors.

B. From occurrences to frequencies

Occurrence count is a good start but there is an issue: longer documents will have higher average count values than shorter documents, even though they might talk about the same topics.

To avoid these potential discrepancies it suffices to divide the number of occurrences of each word in a document by the total number of words in the document: these new features are called **TF** for **Term Frequencies**.

Another refinement on top of TF is to downscale weights for words that occur in many documents in the corpus and are therefore less informative than those that occur only in a smaller portion of the corpus.

This downscaling is called TF-IDF for “**Term Frequency times Inverse Document Frequency**”.

Both **TF** and **TF-IDF** can be computed as follows using TfidfTransformer

Note : You can use another type of transformation, i.e. word2vec, document2vec,...

2. Analysis of the dataset

In order to analyze the dataset, you have to extract some statistical information from the given dataset, for example: the type of data, the missing values, outliers, the correlation between variables, etc. This part should contain also the analysis of the domain application and explain the goal of the analysis and prediction.

3. Machine Learning

1. Use the Regression Line to predict values for some variables (you can choose any variable as a target. Compute the error.
2. Use a machine learning method (clustering and/or classification) in order to predict the class of a new set of objects. You can use the methods as K-Nearest Neighbours (K-NN), Support Vector Machine (SVM), Decision trees, Neural Networks ... The obtained results should be validated using some external indexes as Prediction Error or others. The obtained results should be analyzed in the report and provide a solution to ameliorate the results.

4. Visualization:

You can visualize the knowledge extracted from the classification/clustering in order to present the results i.e. scatter plots using predicted colors, decision trees,...

5. Theoretical details:

Give the algorithmically (mathematical) formalism of the method which give the best results. Explain all the parameters of the used method and their impact on the results.

Some comparison should be made to conclude the project.

The report should contain all these 5 parts, and you should explain the motivation of the used methods, and discuss the obtained results.

Please note that you can use Python, R and/or Matlab and/or Knime for the experimentations.